

# CSE 150

# Operating Systems

Concurrency

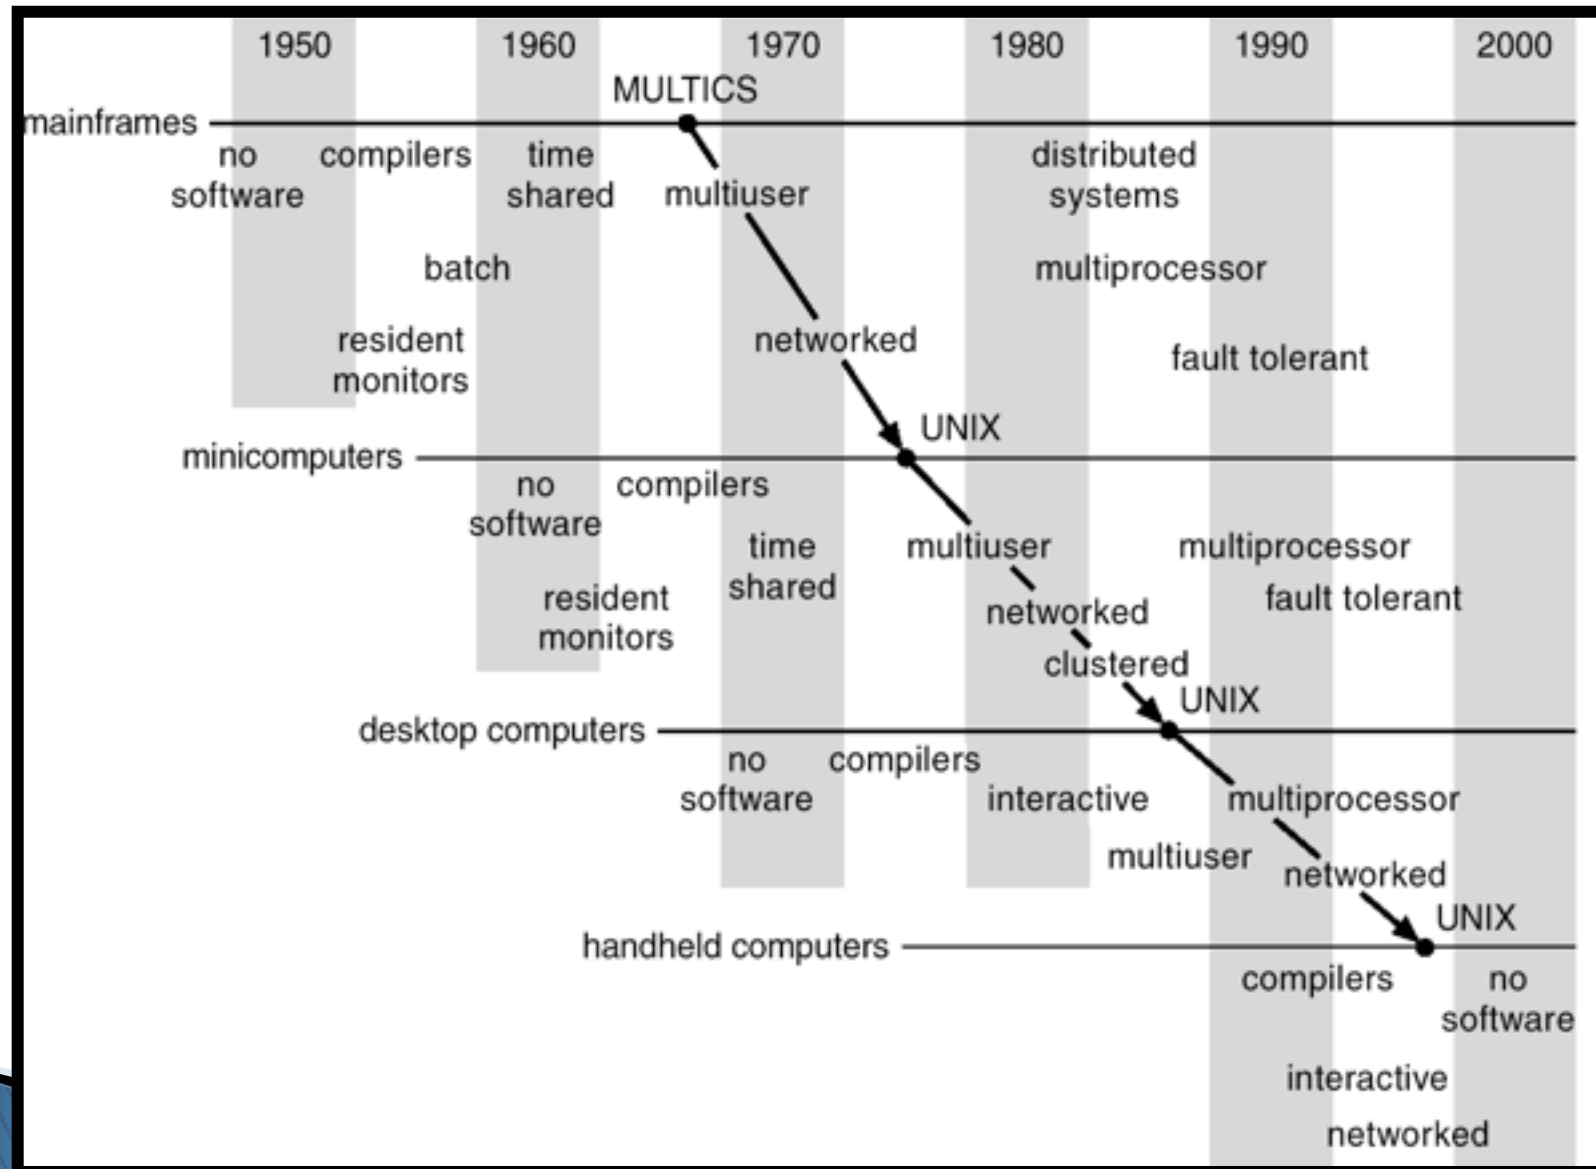


UCMERCED

# Review: History of OS

- ▶ Why Study?
  - To understand how user needs and hardware constraints influenced (and will influence) operating systems
- ▶ Several Distinct Phases:
  - Hardware Expensive, Humans Cheap
    - Eniac, ... Multics
  - Hardware Cheaper, Humans Expensive
    - PCs, Workstations, Rise of GUIs
  - Hardware Really Cheap, Humans Really Expensive
    - Ubiquitous devices, Widespread networking
- ▶ Rapid Change in Hardware Leads to changing OS
  - Batch ⇒ Multiprogramming ⇒ Timeshare ⇒ Graphical UI ⇒ Ubiquitous Devices ⇒ Cyberspace/Metaverse/??
  - Gradual Migration of Features into Smaller Machines
- ▶ Situation today is much like the late 60s
  - Small OS: 100K lines/Large: 10M lines (5M browser!)
  - 100–1000 people-years

# Review: Migration of OS Concepts and Features



# Implementation Issues

## (How is the OS implemented?)

- ▶ Policy vs. Mechanism
  - Policy: **What** do you want to do?
  - Mechanism: **How** are you going to do it?
  - Should be separated, since policies change
- ▶ Algorithms used
  - Linear, Tree-based, Log Structured, etc...
- ▶ Event models used
  - threads vs event loops
- ▶ Backward compatibility issues
  - Very important for Windows 2000/XP/Vista/...
  - POSIX tries to help here
- ▶ System generation/configuration
  - How to make generic OS fit on specific hardware

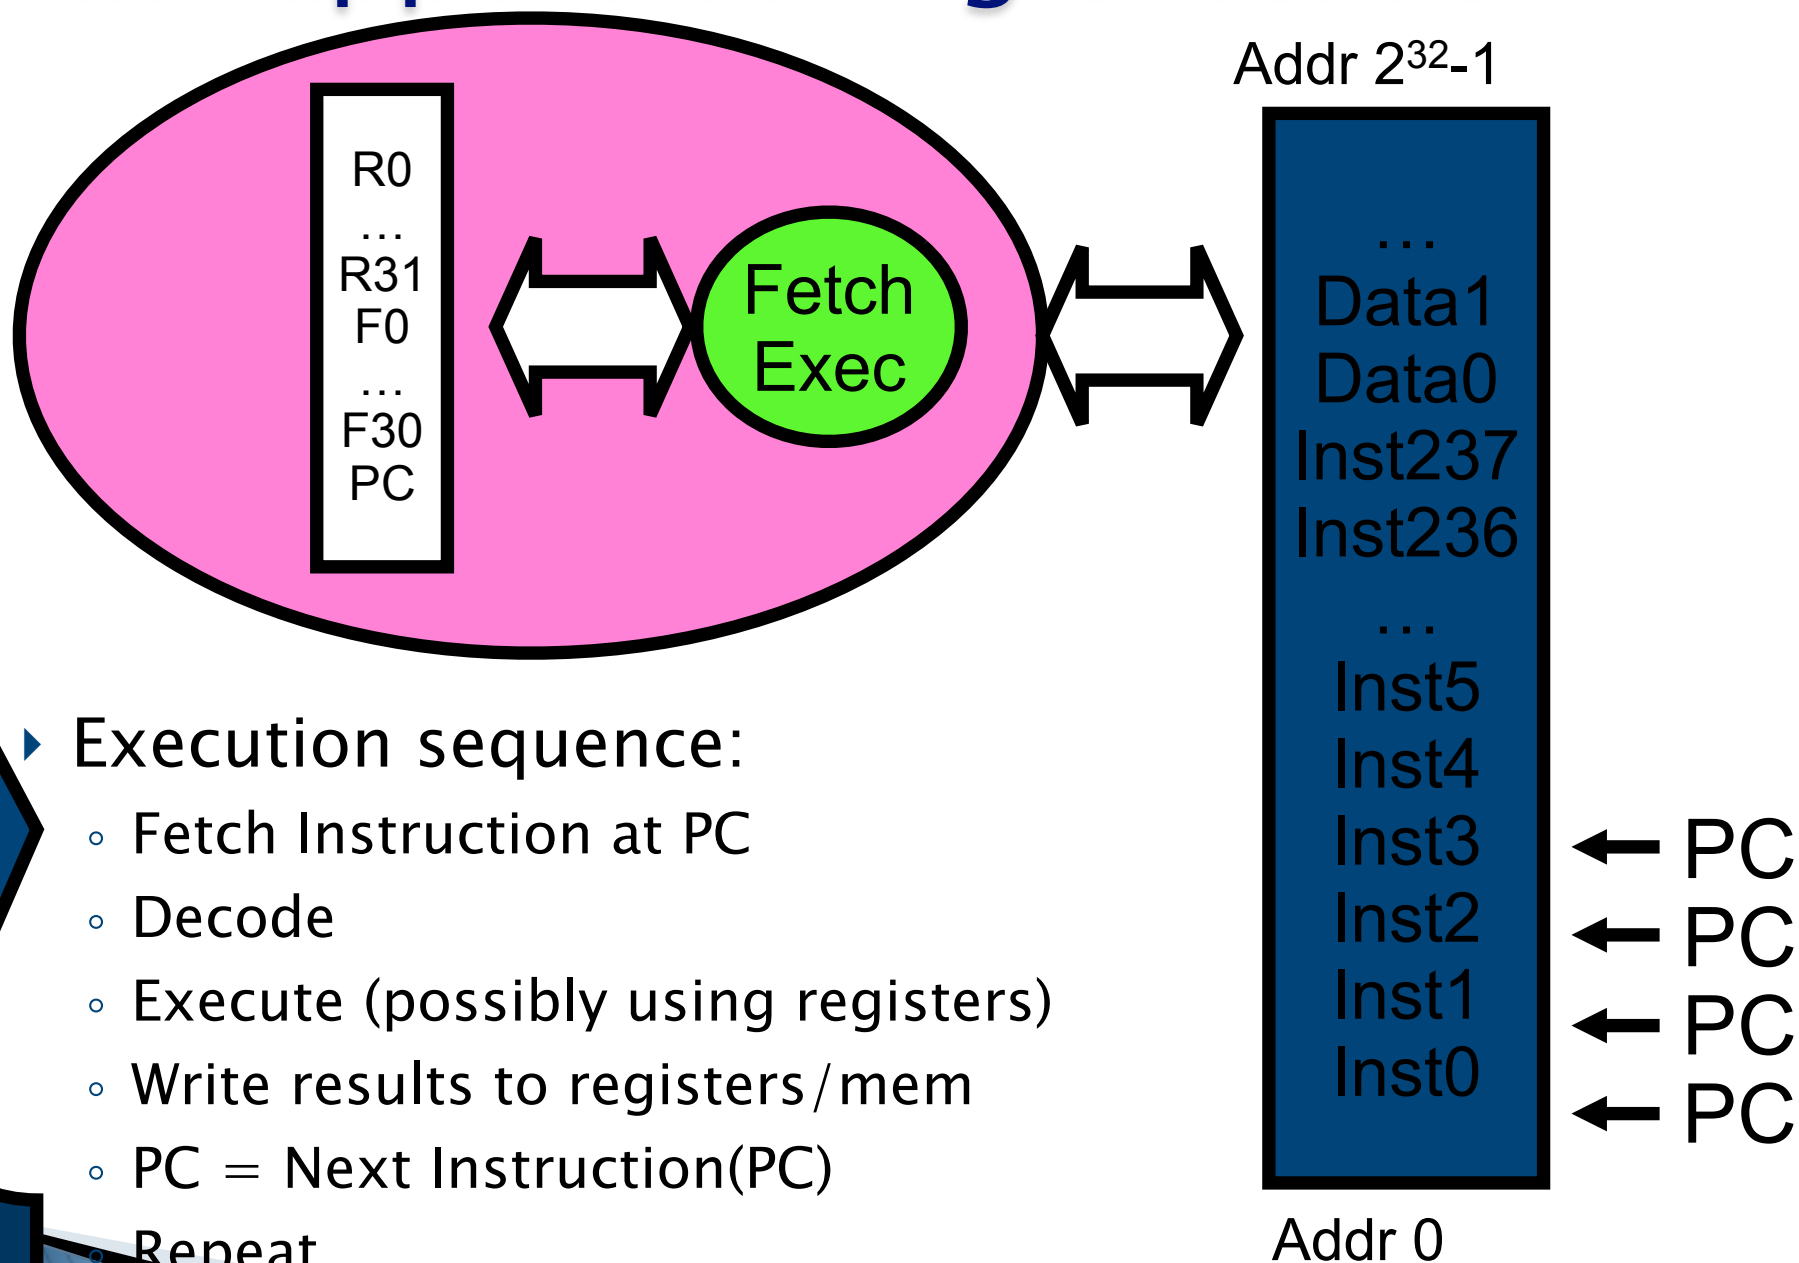
# Concurrency

- ▶ “Thread” of execution
  - Independent Fetch/Decode/Execute loop
  - Operating in some Address space
- ▶ Uniprogramming: one thread at a time
  - MS/DOS, early Macintosh, Batch processing
  - Easier for operating system builder
  - Get rid concurrency by defining it away
  - Does this make sense for personal computers?
- ▶ Multiprogramming: more than one thread at a time
  - Multics, UNIX/Linux, OS/2, Windows NT/2000/XP, Mac OS X
  - Often called “multitasking”, but multitasking has other meanings (talk about this later)
- ▶ ManyCore  $\Rightarrow$  Multiprogramming, right?


# The Basic Problem of Concurrency

- ▶ The basic problem of concurrency involves resources:
  - Hardware: single CPU, single DRAM, single I/O devices
  - Multiprogramming API: users think they have exclusive access to shared resources
- ▶ OS Has to coordinate all activity
  - Multiple users, I/O interrupts, ...
  - How can it keep all these things straight?
- ▶ Basic Idea: Use Virtual Machine abstraction
  - Decompose hard problem into simpler ones
  - Abstract the notion of an executing program
  - Then, worry about multiplexing these abstract machines
- ▶ Dijkstra did this for the “THE system”
  - Few thousand lines vs 1 million lines in OS 360 (1K bugs)

# What happens during execution?



# How can we give the illusion of multiple processors?



- ▶ Assume a single processor. How do we provide the illusion of multiple processors?
  - Multiplex in time!
- ▶ Each virtual “CPU” needs a structure to hold:
  - Program Counter (PC), Stack Pointer (SP)
  - Registers (Integer, Floating point, others...?)
- ▶ How switch from one CPU to the next?
  - Save PC, SP, and registers in current state block
  - Load PC, SP, and registers from new state block
- ▶ What triggers switch?
  - Timer, voluntary yield, I/O, other things



# Properties of this simple multiprogramming technique

- ▶ All virtual CPUs share same non-CPU resources
  - I/O devices the same
  - Memory the same
- ▶ Consequence of sharing:
  - Each thread can access the data of every other thread (good for sharing, bad for protection)
  - Threads can share instructions (good for sharing, bad for protection)
  - Can threads overwrite OS functions?
- ▶ This (unprotected) model common in:
  - Embedded applications
  - Windows 3.1 / Machintosh (switch only with yield)
  - Windows 95—ME? (switch with both yield and timer)

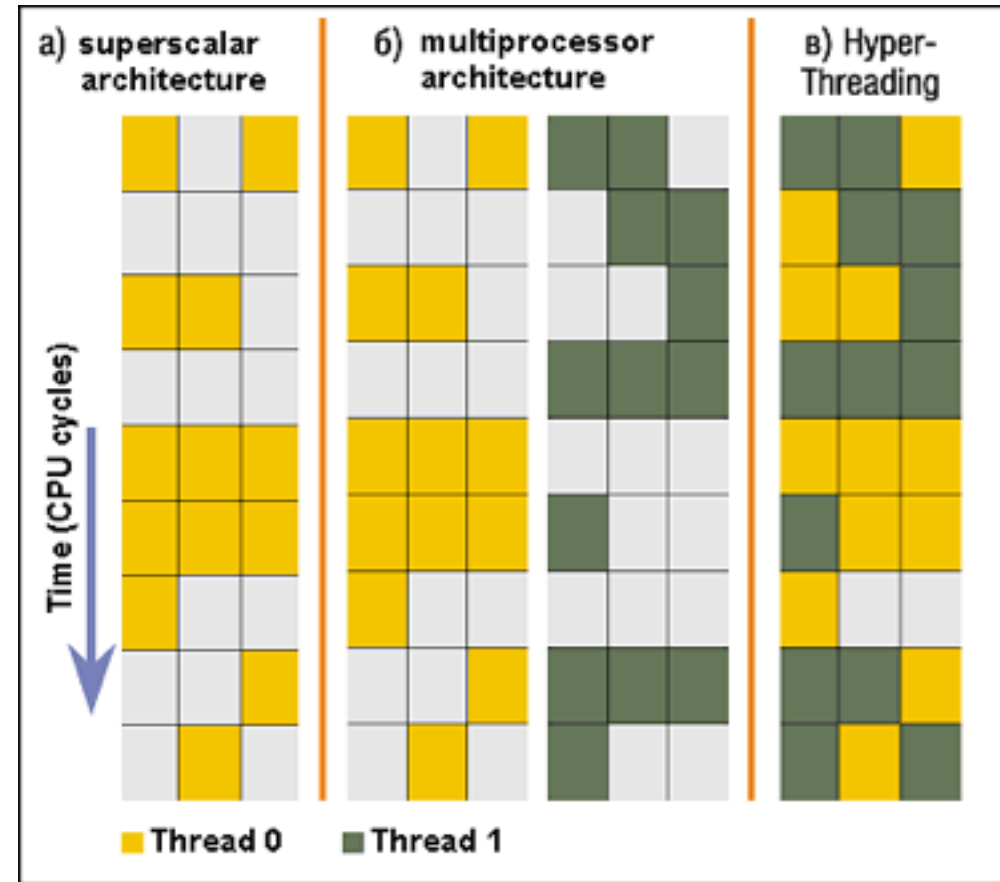
# Modern Technique: SMT/Hyperthreading

- ▶ Hardware technique

- Exploit natural properties of superscalar processors to provide illusion of multiple processors
- Higher utilization of processor resources

- ▶ Can schedule each thread as if were separate CPU

- However, not linear speedup!
- If have multiprocessor, should schedule each processor first



- ▶ Original technique called “Simultaneous Multithreading”

- See <http://www.cs.washington.edu/research/smt/>
- Alpha, SPARC, Pentium 4 (“Hyperthreading”), Power 5

# How to protect threads from one another?

## ► Need three important things:

### 1. Protection of memory

- Every task does not have access to all memory

### 2. Protection of I/O devices

- Every task does not have access to every device


### 3. Protection of Access to Processor:

Preemptive switching from task to task

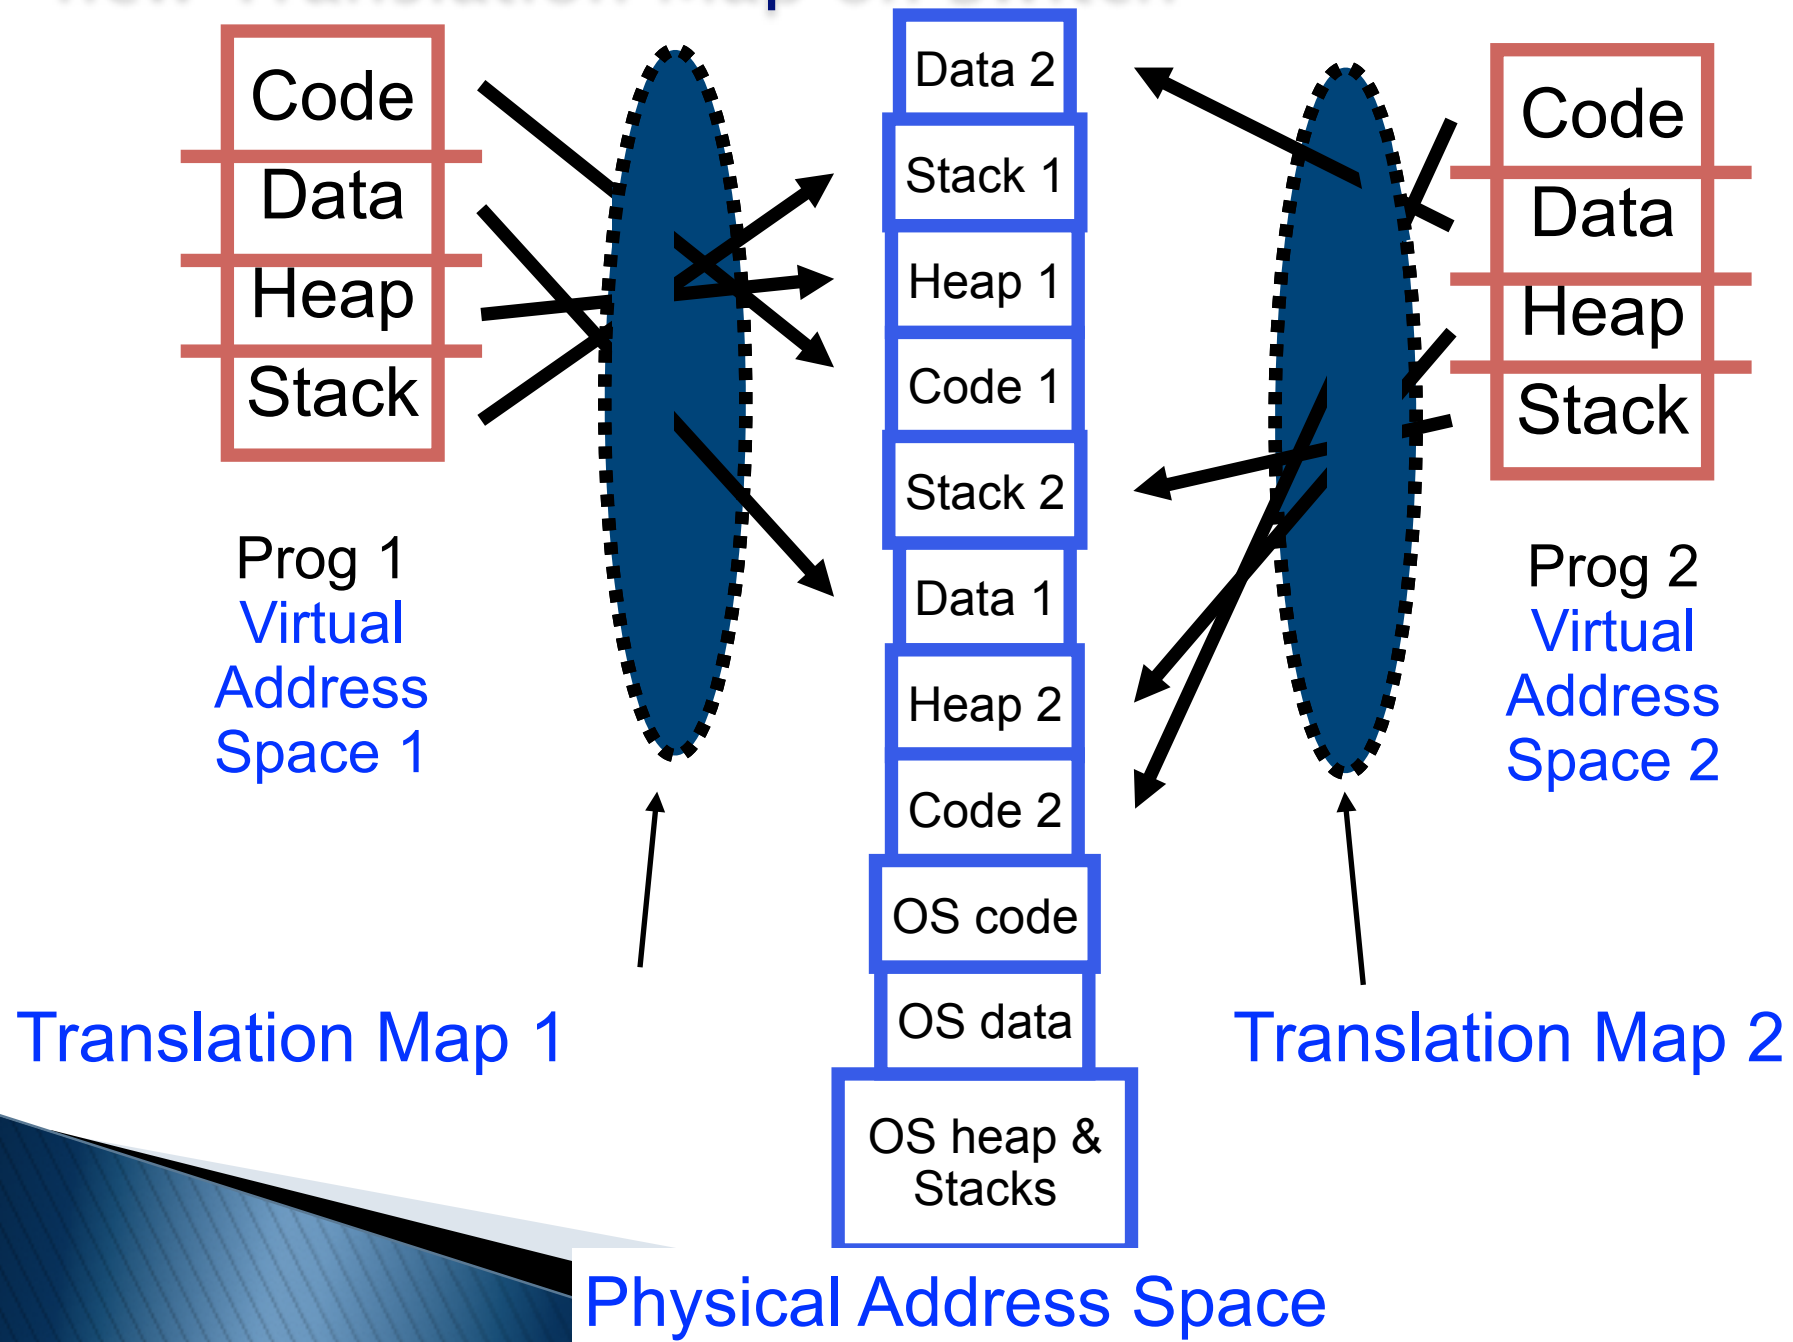
- Use of timer
- Must not be possible to disable timer from usercode

# Recall: Program's Address Space

- ▶ Address space  $\Rightarrow$  the set of accessible addresses + state associated with them:
  - For a 32-bit processor there are  $2^{32} = 4$  billion addresses
- ▶ What happens when you read or write to an address?
  - Perhaps Nothing
  - Perhaps acts like regular memory
  - Perhaps ignores writes
  - Perhaps causes I/O operation
    - (Memory-mapped I/O)
  - Perhaps causes exception (fault)



# Providing Illusion of Separate Address Space: Load new Translation Map on Switch

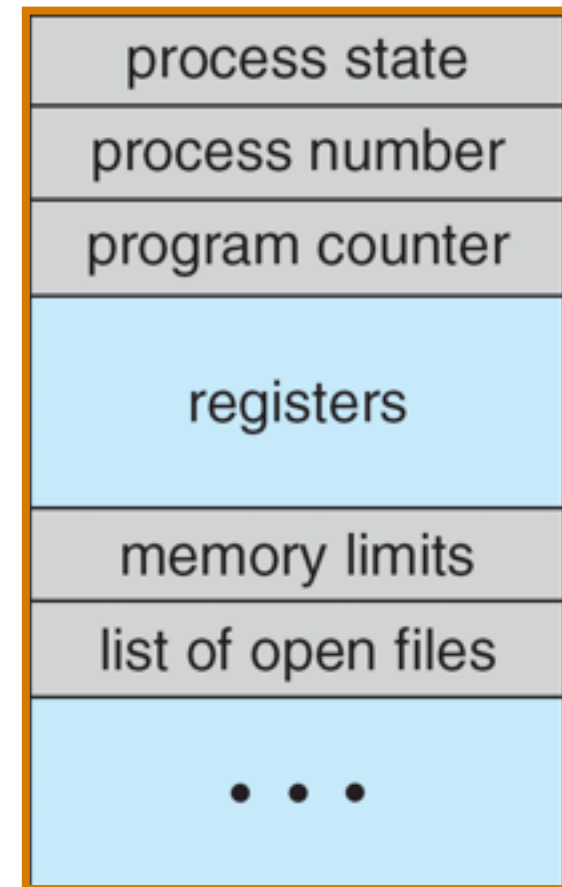


# Traditional UNIX Process

- ▶ Process: Operating system abstraction to represent what is needed to run a single program
  - Often called a “HeavyWeight Process”
  - Formally: a single, sequential stream of execution in its own address space
- ▶ Two parts:
  - Sequential Program Execution Stream
    - Code executed as a single, sequential stream of execution
    - Includes State of CPU registers
  - Protected Resources:
    - Main Memory State (contents of Address Space)
    - I/O state (i.e. file descriptors)
- ▶ Important: There is no concurrency in a heavyweight process

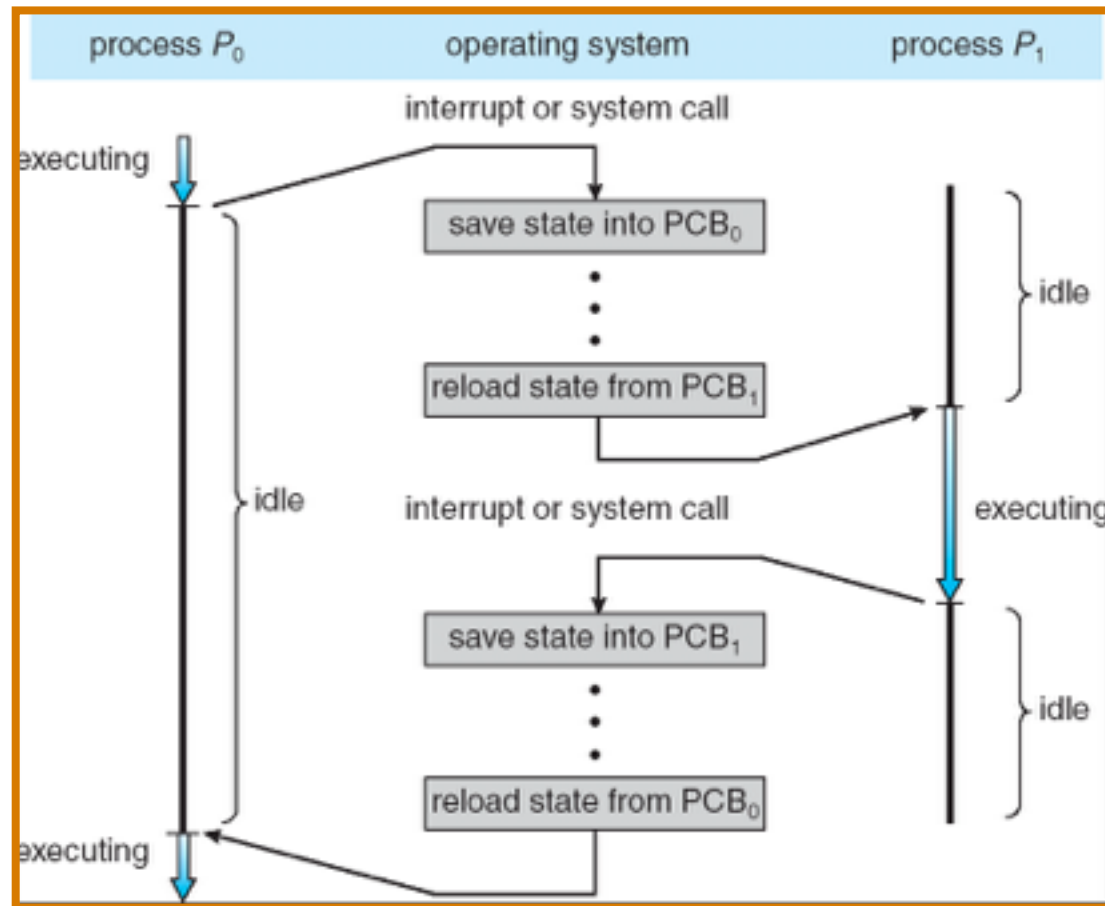
# How do we multiplex processes?

- ▶ The current state of process held in a process control block (PCB):
  - This is a “snapshot” of the execution and protection environment
  - Only one PCB active at a time
- ▶ Give out CPU time to different processes (**Scheduling**):
  - Only one process “running” at a time
  - Give more time to important processes
- ▶ Give pieces of resources to different processes (**Protection**):
  - Controlled access to non-CPU resources
  - Sample mechanisms:
    - Memory Mapping: Give each process their own address space
    - Kernel/User duality: Arbitrary multiplexing of I/O through system calls



Process  
Control  
Block

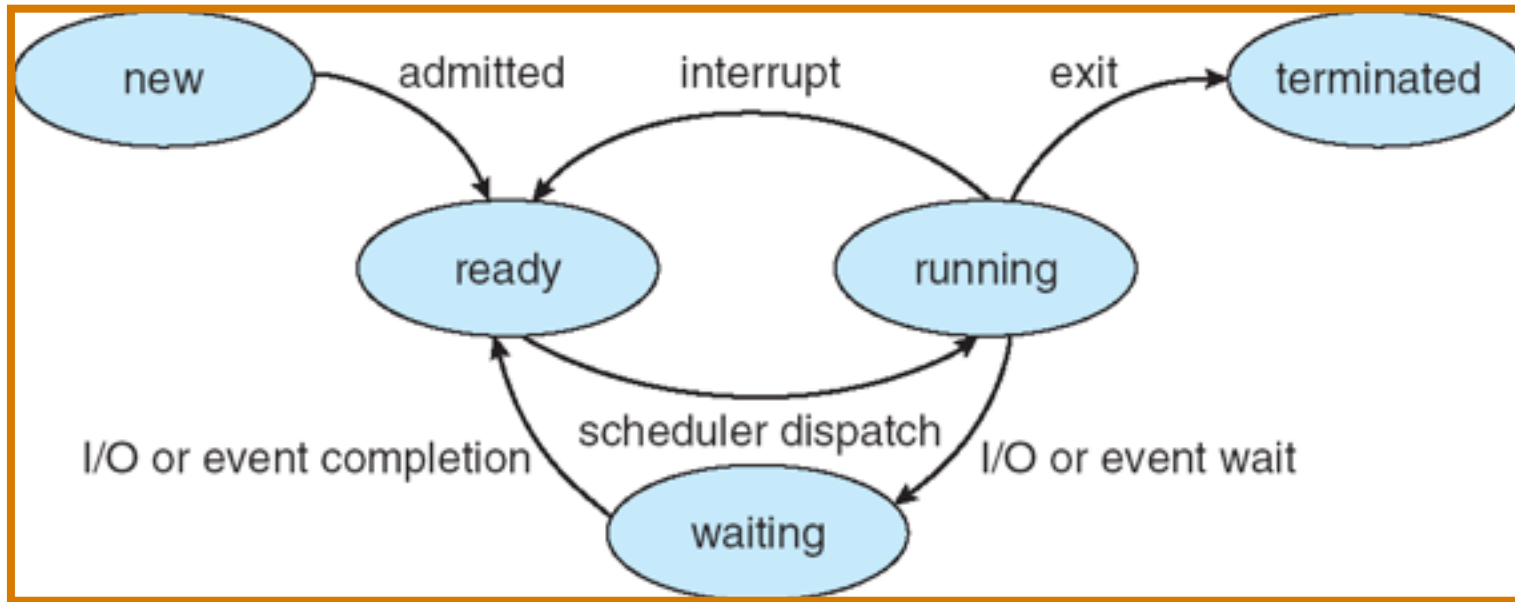
# CPU Switch Processes



- ▶ This is also called a “context switch”
- ▶ Code executed in kernel above is overhead
  - Overhead sets minimum practical switching time
  - Less overhead with SMT/hyperthreading, but... contention for resources instead

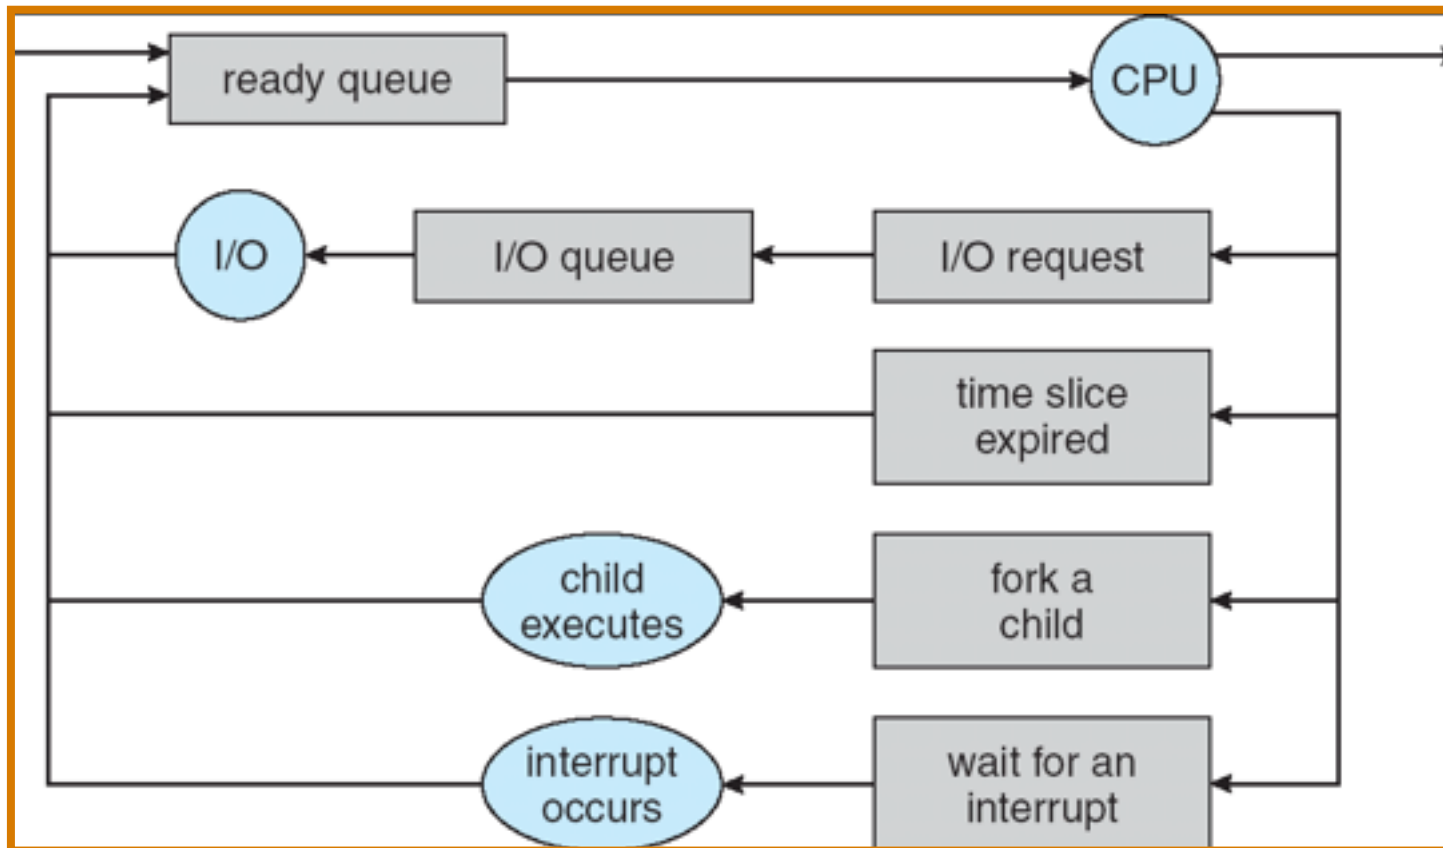


# Diagram of Process State



- ▶ As a process executes, it changes state
  - **new**: The process is being created
  - **ready**: The process is waiting to run
  - **running**: Instructions are being executed
  - **waiting**: Process waiting for some event to occur
  - **terminated**: The process has finished execution

# Process Scheduling



- ▶ PCBs move from queue to queue as they change state
  - Decisions about which order to remove from queues are **Scheduling** decisions
  - Many algorithms possible (few weeks from now)

# What does it take to create a process?

- ▶ Must construct new PCB
  - Inexpensive
- ▶ Must set up new page tables for address space
  - More expensive
- ▶ Copy data from parent process? (Unix `fork()` )
  - Semantics of Unix `fork()` are that the child process gets a complete copy of the parent memory and I/O state
  - Originally very expensive
  - Much less expensive with “copy on write”
- ▶ Copy I/O state (file handles, etc)
  - Medium expense

# Process =? Program

```
main () {  
    ...;  
}  
A () {  
    ...  
}
```

Program

```
main () {  
    ...;  
}  
A () {  
    ...  
}
```


Heap

Stack  
A  
main

Process


- ▶ More to a process than just a program:
  - Program is just part of the process state
  - I run emacs on lectures.txt, you run it on homework.java – Same program, different processes
- ▶ Less to a process than a program:
  - A program can invoke more than one process
  - cc starts up cpp, cc1, cc2, as, and ld

# Multiple Processes Collaborate on a Task



- ▶ High Creation/memory Overhead
- ▶ (Relatively) High Context-Switch Overhead
- ▶ Need Communication mechanism:
  - Separate Address Spaces Isolates Processes
  - Shared-Memory Mapping
    - Accomplished by mapping addresses to common DRAM
    - Read and Write through memory
  - Message Passing
    - `send()` and `receive()` messages
    - Works across network

# Shared Memory Communication



- ▶ Communication occurs by “simply” reading/writing to shared address page
  - Really low overhead communication
  - Introduces complex synchronization problems

# Inter-process Communication (IPC)

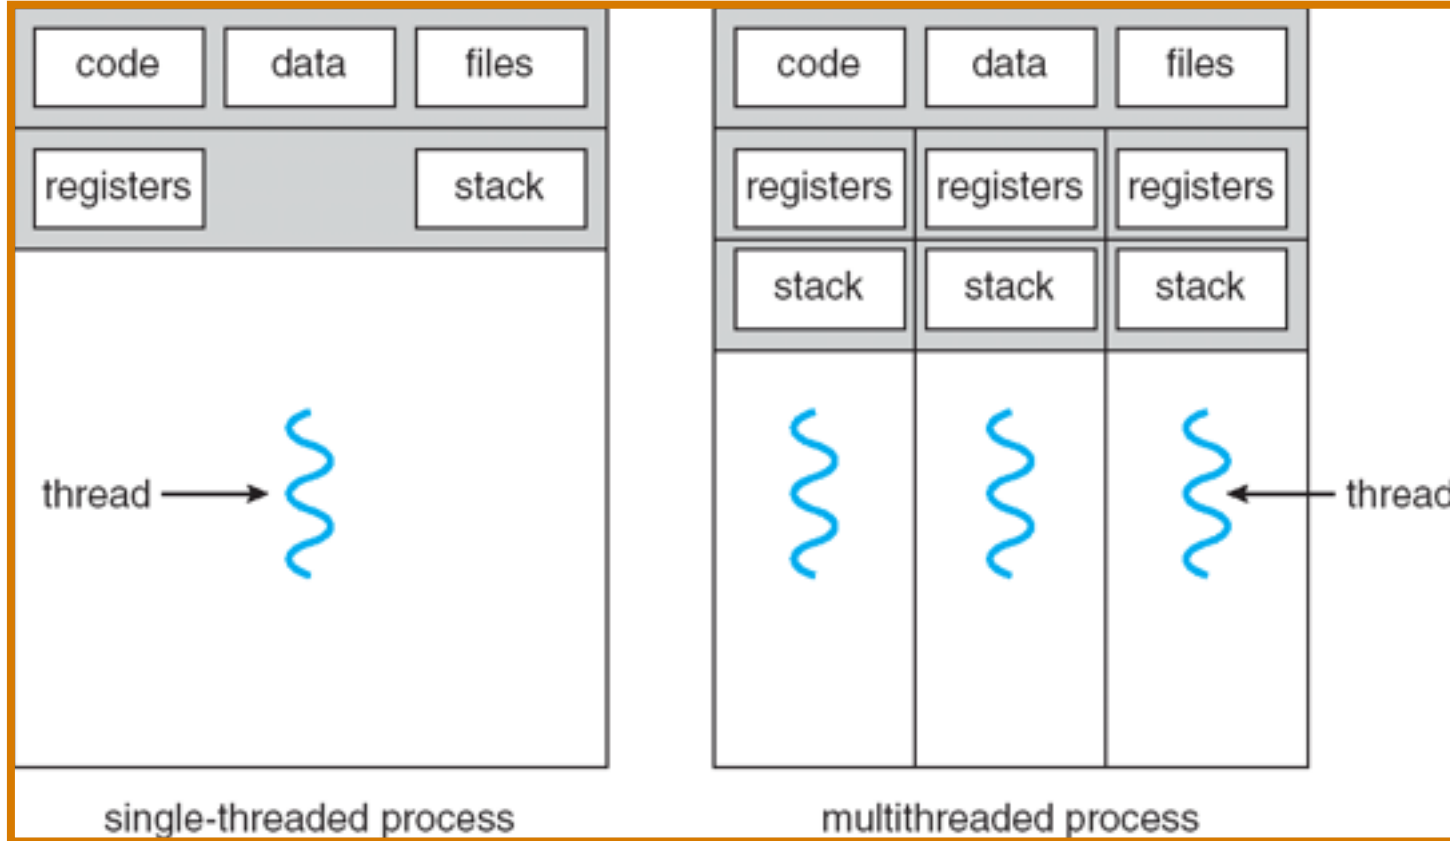
- ▶ Mechanism for processes to communicate and to synchronize their actions
- ▶ Message system – processes communicate with each other without resorting to shared variables
- ▶ IPC facility provides two operations:
  - `send(message)` – message size fixed or variable
  - `receive(message)`
- ▶ If P and Q wish to communicate, they need to:
  - establish a communication link between them
  - exchange messages via send/receive
- ▶ Implementation of communication link
  - physical (e.g., shared memory, hardware bus, syscall/trap)
  - logical (e.g., logical properties)

# Modern “Lightweight” Process with Threads

- ▶ Thread: a sequential execution stream within process (Sometimes called a “Lightweight process”)
  - Process still contains a single Address Space
  - No protection between threads
- ▶ Multithreading: a single program made up of a number of different concurrent activities
  - Sometimes called multitasking, as in Ada...
- ▶ Why separate the concept of a thread from that of a process?
  - Discuss the “thread” part of a process (concurrency)
  - Separate from the “address space” (Protection)
  - Heavyweight Process  $\equiv$  Process with one thread



# Single and Multithreaded Processes



- ▶ Threads encapsulate concurrency: “Active” component
- ▶ Address spaces encapsulate protection: “Passive” part
  - Keeps buggy program from trashing the system
- ▶ Why have multiple threads per address space?

# Examples of multithreaded programs

## ▶ Embedded systems

- Elevators, Planes, Medical systems, Wristwatches
- Single Program, concurrent operations

## ▶ Most modern OS kernels

- Internally concurrent because have to deal with concurrent requests by multiple users
- But no protection needed within kernel

## ▶ Database Servers

- Access to shared data by many concurrent users
- Also background utility processing must be done

# Examples of multithreaded programs (con't)

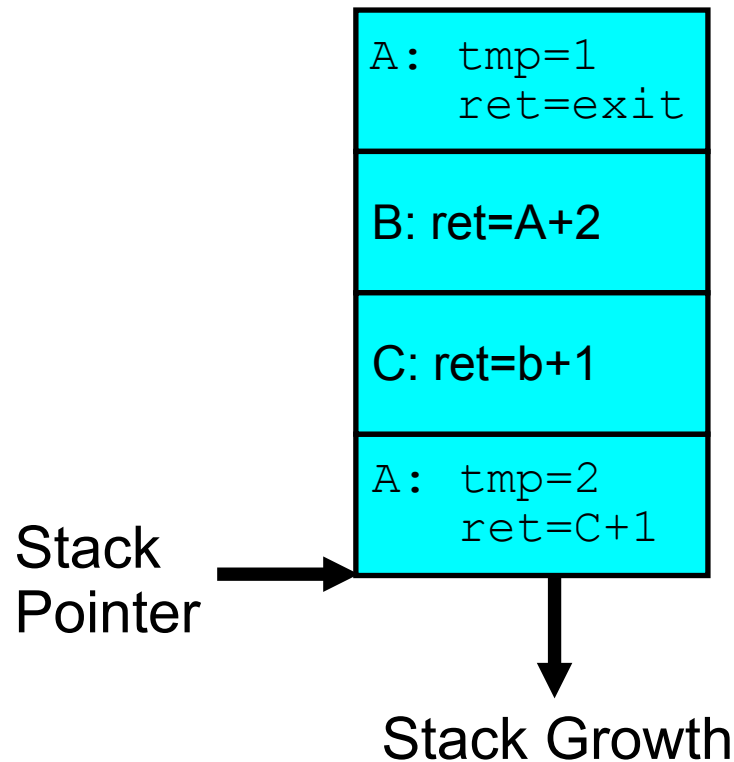
- ▶ Network Servers
  - Concurrent requests from network
  - Again, single program, multiple concurrent operations
  - File server, Web server, and airline reservation systems
- ▶ Parallel Programming (More than one physical CPU)
  - Split program into multiple threads for parallelism
  - This is called Multiprocessing
- ▶ Some multiprocessors are actually uniprogrammed:
  - Multiple threads in one address space but one program at a time

# Thread State

- ▶ State shared by all threads in process/addr space
  - Contents of memory (global variables, heap)
  - I/O state (file system, network connections, etc)
- ▶ State “private” to each thread
  - Kept in TCB  $\equiv$  Thread Control Block
  - CPU registers (including, program counter)
  - Execution stack – what is this?
- ▶ Execution Stack
  - Parameters, Temporary variables
  - return PCs are kept while called procedures are executing

# Execution Stack Example

```
A(int tmp) {  
    if (tmp<2)  
        B();  
    printf(tmp);  
}  
B() {  
    C();  
}  
C() {  
    A(2);  
}  
A(1);
```



- ▶ Stack holds temporary results
- ▶ Permits recursive execution
- ▶ Crucial to modern languages

# Classification


# threads Per AS:	# of addr spaces:	One	Many
One		MS/DOS, early Macintosh	Traditional UNIX
Many		Embedded systems (Geoworks, VxWorks, JavaOS, etc) JavaOS, Pilot(PC)	Mach, OS/2, Linux Windows 9x??? Win NT to XP, Solaris, HP-UX, OS X

- ▶ Real operating systems have either
  - One or many address spaces
  - One or many threads per address space
- ▶ Did Windows 95/98/ME have real memory protection?
  - No: Users could overwrite process tables/System DLLs

# Example: Implementation Java OS

- ▶ Many threads, one Address Space
- ▶ Why another OS?
  - Recommended Minimum memory sizes:
    - UNIX + X Windows: 32MB
    - Windows 98: 16–32MB
    - Windows NT: 32–64MB
    - Windows 2000/XP: 64–128MB
  - What if we want a cheap network point-of-sale computer?
    - Say need 1000 terminals
    - Want < 8MB
- ▶ What language to write this OS in?
  - C/C++/ASM? Not terribly high-level. Hard to debug.
  - Java/Lisp? Not quite sufficient – need direct access to HW/memory management

## Java OS Structure



# Summary

- ▶ Processes have two parts
  - Threads (Concurrency)
  - Address Spaces (Protection)
- ▶ Concurrency accomplished by multiplexing CPU Time:
  - Unloading current thread (PC, registers)
  - Loading new thread (PC, registers)
  - Such context switching may be voluntary (`yield()`, I/O operations) or involuntary (timer, other interrupts)
- ▶ Protection accomplished restricting access:
  - Memory mapping isolates processes from each other
  - Dual-mode for isolating I/O, other resources
- ▶ Book talks about processes
  - When this concerns concurrency, really talking about thread portion of a process
  - When this concerns protection, talking about address space portion of a process